

6.S081: Introduction

Adam Belay
abelay@mit.edu

6.S081 Objectives

- Understand how OSes are designed and implemented
- Hands-on experience building systems software
- Will extend a simple OS (xv6)
- Will learn about how hardware works (Risc-V)

Some things you'll do in 6.S081

1. You will build a driver for a network stack that sends packets over the real Internet
2. You will redesign a memory allocator so that it can scale across multiple cores
3. You will implement fork and make it efficient through an optimization called copy-on-write

What is the purpose of an OS?

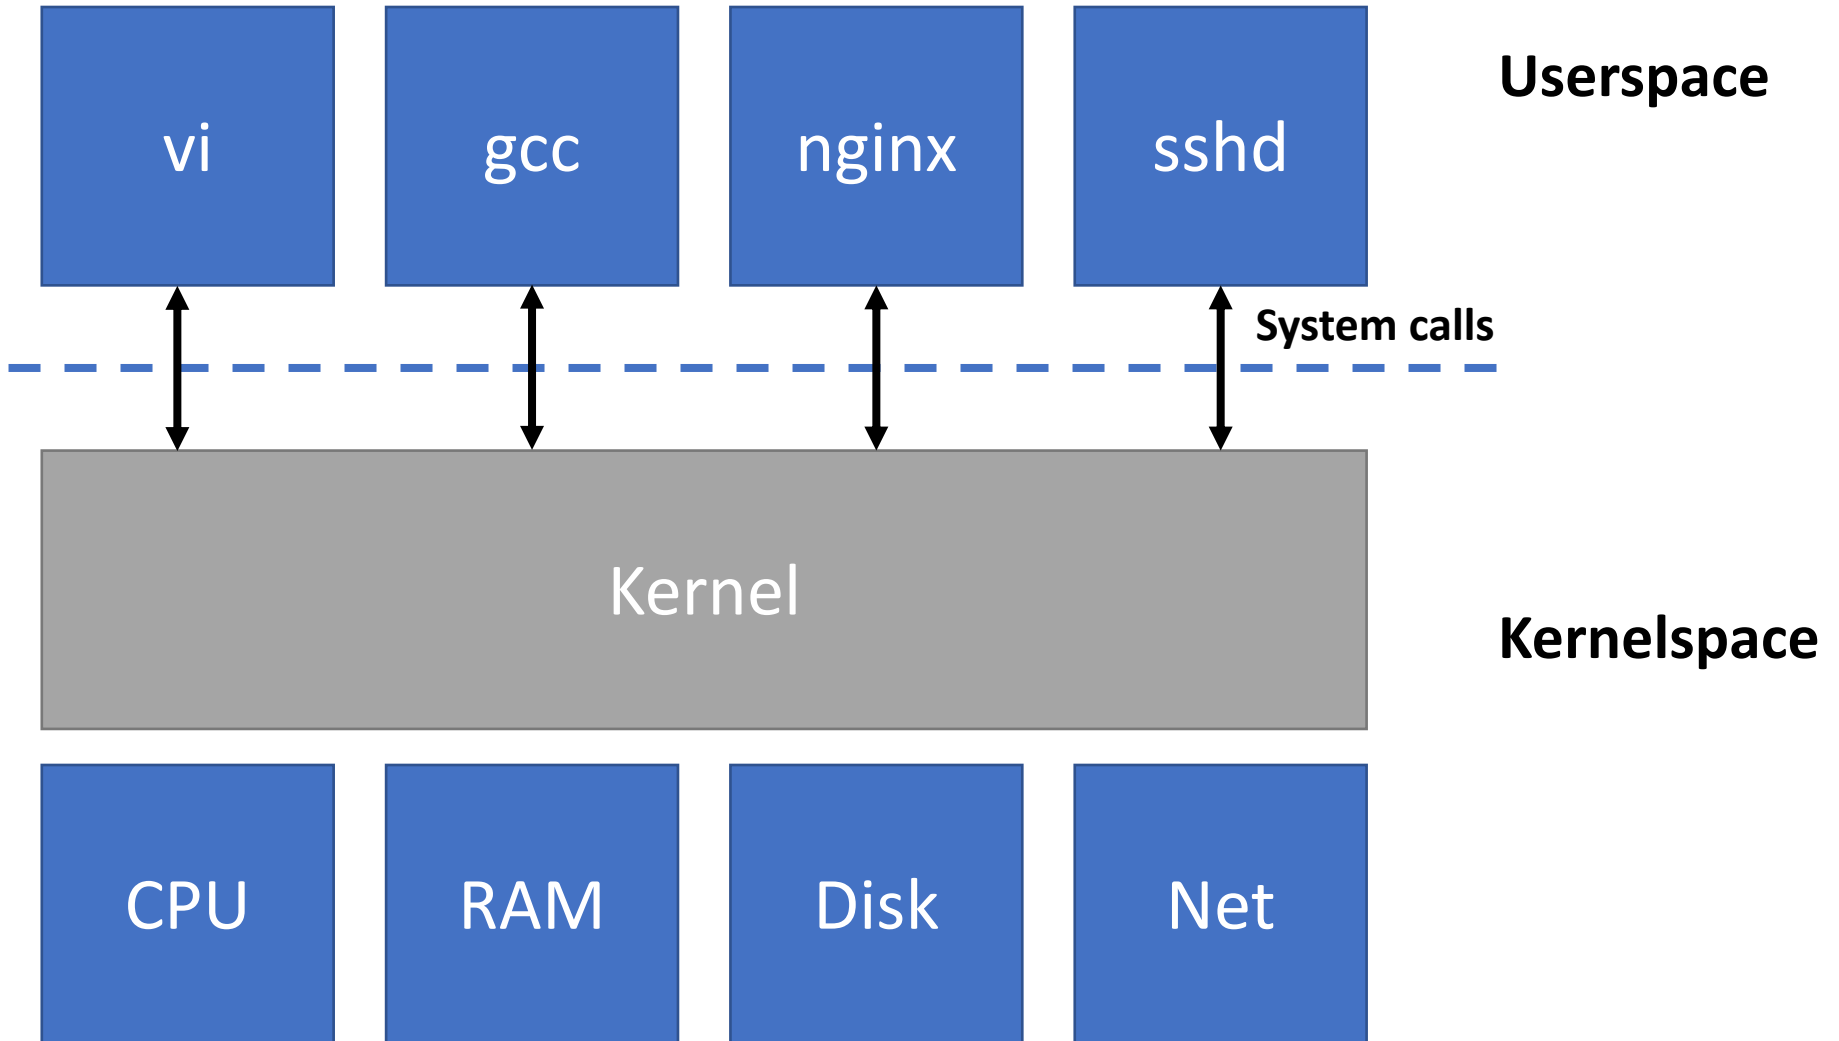
1. Abstraction

- Hides hardware details for portability and convenience
- Must not get in the way of high performance
- Must support a wide range of applications

2. Multiplexing

- Allows multiple applications to share hardware
- Isolation to contain bugs and provide security
- Sharing to allow cooperation

OS Organization



OS abstractions

- Process (a running program)
- Memory allocation
- File descriptors
- File names and directories
- Access control and quotas
- Many others: users, IPC, network sockets, time, etc.

User <-> kernel interface

- Primarily system calls

- Examples:

```
fd = open("out", 1);
```

```
len = write(fd, "hello\n", 6);
```

```
pid = fork();
```

- Look and behave like function calls, but they aren't

Why OSes are interesting

- Unforgiving to build: Debugging is hard, a single bug can take down the entire machine
- Design tensions:
 - Efficiency vs. Portability/Generality
 - Powerful vs. Simple
 - Flexible vs. Secure
- Challenge: good orthogonality, feature interactions
- Varied uses from smartbulbs to supercomputers
- Evolving HW: NVRAM, Multicore, 200Gbit networks

Take this course if you:

- Want to understand how computers really work from an engineering perspective
- Want to build future system infrastructure
- Want to solve bugs and security problems
- Care about performance

Logistics

Online resources

- Course website
 - <https://pdos.csail.mit.edu/6.S081/>
 - Schedule, course policies, lab assignments, etc.
 - Videos and notes of 2020 lectures
- Piazza
 - <https://piazza.com/mit/fall2021/6s081>
 - Announcements and discussion
 - Ask questions about labs and lecture

Lectures

1. OS concepts
 2. Case studies of xv6 --- a simple, small OS
 3. Lab background and solutions
 4. OS papers
- Submit a question before each lecture
 - Resource: xv6 book

Labs

- Goal: Hands-on experience
- Three types of labs:
 1. Systems programming: due next week
 2. OS primitives: e.g., thread scheduling
 3. OS extensions: e.g., networking driver

Collaboration

- Feel free to ask and discussion questions about lab assignments in class or on Piazza
- Discussion is great
 - But all solutions (code and written work) must be your own
 - Acknowledge ideas from others (e.g., classmates, open source software, stackoverflow, etc.)
- Do not post your solutions (including on github)

Covid-19 and in-person learning

- Masks are **required**; must be worn correctly
- If you have symptoms or test positive...
 - Don't attend class, contact us right away
 - We will work with you to provide course materials

Grading

- 70% labs, based on the same tests you will run
- 20% lab check off meetings
 - We will ask questions about randomly selected labs during office hours
- 10% homework and class/piazza participation

Back to system calls

- I'll show examples of using system calls
- Will use xv6, the same OS you'll build labs on
- xv6 is similar to UNIX or Linux, but way simpler
 - Why? So you can understand the entire thing.
- Why UNIX?
 - Clean design, widely used: Linux, OSX, Windows (mostly)
- xv6 runs on Risc-V, like 6.004
- You will use Qemu to run xv6 (emulation)

System call

Description

<code>int fork()</code>	Create a process, return child's PID.
<code>int exit(int status)</code>	Terminate the current process; status reported to <code>wait()</code> . No return.
<code>int wait(int *status)</code>	Wait for a child to exit; exit status in <code>*status</code> ; returns child PID.
<code>int kill(int pid)</code>	Terminate process PID. Returns 0, or -1 for error.
<code>int getpid()</code>	Return the current process's PID.
<code>int sleep(int n)</code>	Pause for n clock ticks.
<code>int exec(char *file, char *argv[])</code>	Load a file and execute it with arguments; only returns if error.
<code>char *sbrk(int n)</code>	Grow process's memory by n bytes. Returns start of new memory.
<code>int open(char *file, int flags)</code>	Open a file; flags indicate read/write; returns an fd (file descriptor).
<code>int write(int fd, char *buf, int n)</code>	Write n bytes from buf to file descriptor fd; returns n.
<code>int read(int fd, char *buf, int n)</code>	Read n bytes into buf; returns number read; or 0 if end of file.
<code>int close(int fd)</code>	Release open file fd.
<code>int dup(int fd)</code>	Return a new file descriptor referring to the same file as fd.
<code>int pipe(int p[])</code>	Create a pipe, put read/write file descriptors in <code>p[0]</code> and <code>p[1]</code> .
<code>int chdir(char *dir)</code>	Change the current directory.
<code>int mkdir(char *dir)</code>	Create a new directory.
<code>int mknod(char *file, int, int)</code>	Create a device file.
<code>int fstat(int fd, struct stat *st)</code>	Place info about an open file into <code>*st</code> .
<code>int stat(char *file, struct stat *st)</code>	Place info about a named file into <code>*st</code> .
<code>int link(char *file1, char *file2)</code>	Create another name (file2) for the file file1.
<code>int unlink(char *file)</code>	Remove a file.